CSE322

Assignment - DVR Implementation

1 Introduction

In this assignment you will implement Distance Vector Routing (DVR) protocol. You must word with **UNIX sockets**, so you have to use a UNIX based machine. The codes were tested on Ubuntu, but pretty much any Linux distro/Mac OS should work just fine. Moreover, as you will be configuring **iptables** and ports, it is strongly suggested that you work on a **Virtual Machine** instead of your native OS. Chances are, even if things go wrong it won't do any permanent damage - but let's not take the risk at the first place. **You have to complete this assignment in C++.**

2 Overview

For this assignment we are going to provide **three** files to you -

1. a setup file (setup),

2. a topology file (topo.txt), and

3. a driver program (driver.py).

Your job is to write the fourth file - the **router program**.

The setup file configures virtual interfaces so that we can emulate routers having different IP addresses on the same machine, the topology file describes the network topology, and the driver program provides necessary commands to the routers. All messages are transmitted among entities through User Datagram Protocol.

3 The Setup File

The setup file contains some **ifconfig** commands. Bascially it creates some virtual network interfaces and attaches an ip address to it. Let's take a closer look at the first command, all other commands are similar:

```
sudo ifconfig eth0:1 192.168.10.1 netmask 255.255.255.0 up
```

Notice that we have to execute these commands with elevated privileges. The first parameter to **ifconfig** is the virtual interface, the second is its IP address. **Keep in mind that based on your OS, interface name might differ**. We will configure **eth0:x** to have the IP address **192.168.10.x**. Finally we are assigning the network mask to be **255.255.255.0** and activating the interface. We are going to assume **driver program's address is 192.168.10.100**.

For the most part of this assignment, you won't have to modify this setup file. The only case when you have to modify this is if we ask you to create a topology with more than four routers. For example, if we ask you to create a topology with 6 routers, you will copy any line and paste it twice (we already have five addresses setup, four routers and one driver, so we need two more routers). Then you are going to change network interface of the pasted lines to eth0:5 and eth0:6, and IP addresses to 192.168.10.5 and 192.168.10.6. Don't change the netmask. Finally you are going to execute this setup file. As a rule of thumb, always execute this setup file everytime you change it.

4   The Topology File

The topology file describes our network topology through a bunch of statements having the following format:

```
<router1 address>     <router2 address>     <link cost>
```

So we get two infomation from each of the statements:

1. there is a link between router1 and router2 i.e. they can send and receive messages to each other, and

2. the cost associated with that link.

For you convenience the contents of the supplied topology file - 'topo.txt' along with the topology it describes is illustrated in Figure 1.

```
192.168.10.1 192.168.10.2 10
192.168.10.2 192.168.10.3 2
192.168.10.1 192.168.10.3 3
192.168.10.2 192.168.10.4 5
192.168.10.3 192.168.10.4 11
```


Figure 1: Network topology represented by topo.txt

Be advised that when we check your offline, we might tell you to create a new topology file or modify the existing one. Depending on how many routers we want in our topology, **you might have to modify the setup file**.

5   The Driver Program

The driver program is the **GOD-entity** - it knows the whole topology, it can update link cost, it can deactivate a link, it can reactivate the deactivated link, it can ask a router to send a message to another router. However, we are going to assume it does not have the authority to create a totally new link. For example, in the topology described previously, there is no link between 192.168.10.1 and 192.168.10.4. The driver will not be able to create a link between these two. But there is a link between 192.168.10.1 and 192.168.10.2 having a cost of 10. The driver has the authority to deactivate this link, reactivate this link, and change the link cost. We can input commands to the driver from the terminal. For your assignment you will not really need to know much about this driver other than the information provided in the following subsections, but feel free to go through the code if it interests you. Be advised, although you are allowed to examine the driver program, **you are not allowed to modify the driver code in any way**.

5.1   Technical Details

- The driver program is written in python 2.x, it has not been tested on python 3.x.

- It binds itself to the IP address 192.168.10.100 and the port 4747.

- It assumes the routers will bind on the same port i.e. 4747.

- You can assume the driver will be started only after all routers are up and running.

- The driver can be invoked by `python driver.py <topology-file>`.

5.2   Available Commands

- **help** - shows a list of available commands.

- **hosts** - shows a list of routers.

- **cost** $\langle ip1 \rangle$ $\langle ip2 \rangle$ $\langle cost \rangle$ - updates link cost to $\langle cost \rangle$ between routers having ip addresses of $\langle ip1 \rangle$ and $\langle ip2 \rangle$. Be advised, this message is sent to only the relevant routers. Once this command is issued, the driver writes a message of the following format to the socket:

  `cost<ip1><ip2><value>`

  Here 'cost' takes up 4 bytes, 'ip1' and 'ip2' take up 4 bytes (one for each of its segment) each, and 'value' takes up 2 bytes. So in total 14 bytes are sent out through the socket.

- **down** $\langle ip1 \rangle$ $\langle ip2 \rangle$ - deactivates link between routers having ip addresses of $\langle ip1 \rangle$ and $\langle ip2 \rangle$. When we 'deactivate' a link, no packets can be sent using that link. Be advised that this change is sustained till you counter it with the following **up** command. It calls a shell command that requires elevated privileges. The restriction is automatically imposed by the OS. Your job is to recognize a link failure and make necessary changes to your routing table. More on this in subsection 6.4.

- **up** $\langle ip1 \rangle$ $\langle ip2 \rangle$ - reactivates link between routers having ip addresses of $\langle ip1 \rangle$ and $\langle ip2 \rangle$. Be advised that this change is sustained till you counter it with the above **down** command. It calls a shell command that requires elevated privileges.

- **send** $\langle ip1 \rangle$ $\langle ip2 \rangle$ $\langle message - length \rangle$ $\langle message \rangle$ - instructs the router having ip address of $\langle ip1 \rangle$ to send $\langle message \rangle$ of length $\langle message - length \rangle$ to the router having ip address of $\langle ip2 \rangle$. Be advised, this message is sent to the first router only. Once this command is issued, the driver writes a message of the following format to the socket:

4

```
send<ip1><ip2><message-length><message>
```

Here 'send' takes up 4 bytes, 'ip1' and 'ip2' take up 4 bytes (one for each of its segment) each, 'message-length' takes up 2 bytes, and 'message' takes up 'message-length' number of bytes. Assuming our message is `hello world!`, 26 bytes are sent through the socket (message being 12 characters long).

- **show** ⟨*ip*⟩ - instructs router having ip address of ⟨*ip*⟩ to print its routing table. Once this command is issued, the driver writes a message of the following format to the socket:

  ```
  show<ip>
  ```

  Here 'show' takes up 4 bytes, 'ip' takes up 4 bytes (one for each of its segment).

- **clear** - clears the screen.

- **exit** - terminates the driver.

5.3 Link 'up' and 'down' Command - What's Happening under the Hood

In subsection 5.2 you were introduced with two driver commands: **up** and **down**. In this subsection we are going to take a look at the 'magic' that's happening behind the scenes. This subsection has been added so that you get a clear idea exactly how we are forcing a particular link to fail. You can safely skip this if you want - it is not going to help you in completing your assignment directly. That being said, a bit of extra knowledge will never harm you.

There is a tool named **iptables** which is used to maintain IPv4 packet filtering rules. Using this very tool we can force packets to be dropped between two IP addresses. Let's take a look at a particular command to get a practical idea:

```
sudo iptables -I OUTPUT -s 192.168.10.1 -d 192.168.10.2 -j DROP
```

The above command basically says: insert a rule in our iptables (iptables -I) specifying that, if we send out (OUTPUT) a packet from source (-s 192.168.10.1) to destination (-d 192.168.10.2), simply drop it (-j DROP). As by link failure we mean no router can send packets to the other, we have to write another command with source and destination interchanged. To change our rule so that the packets are accepted, we need to invoke the above command again with a minor change - we need to replace 'DROP' with 'ACCEPT'.

That was a pretty shallow intro to **iptables**, if this seems interesting to you, simply check out the manpage for more details.

6   Your Task - the Router Program

You have to write the router in **C++**. Suppose you compile your **router.cpp** file into the executable **router**. Assuming you are on the same directory as your executable, you should be able to invoke your router as

```
./router <ip address> <topology-file>
```

A step by step functionality of a router is presented in the following subsections.

6.1   Host and Neighbor List Formation

A router should extract list of routers in the network from the topology file. It should be able to find out its neighbors from the topology file along with associated links. Be advised it should not store any other information from the topology file. For example, in our given topology 192.168.10.1 is connected to both 192.168.10.2 and 192.168.10.3. So the cost of these two links should be stored by our first router, but information regarding all other links should be ignored. Once a router is up and running, it should print its initial routing table once.

6.2   Routing Table Maintenance

Each router should maintain its own routing table. The routing table should contain entries for each of the routers in the network. It should store two information in each entry - next hop to reach the destination and the total cost to reach the destination. For example, first time router 192.168.10.1 is run, its routing table will be something like this:

```
destination      next hop        cost
-----------      --------        ----
192.168.10.2     192.168.10.2    10
192.168.10.3     192.168.10.3    3
192.168.10.4         -           <some-predefined-value>
```

The last entry denotes that the final destination is unreachable. The routing table might need to be updated on the following occasions:

1. when a neighbor sends its routing table,

2. when a link cost is updated,

3. when a link is deactivated, and

4. when a link is reactivated.

You are going to print the routing table once at startup, and whenever `show <ip>` command is invoked.
Bonus: You are recommended to utilize both Split Horizon Rule and Forced Update.

6.3 Periodic Update of Routing Table

Each router will exchange its routing table with its neighbors only. The exchange should be periodical at five seconds interval.

6.4 Detection of Link Failure

We can deactivate a link by invoking the **down** command from our driver, as described in subsection 5.2. A link will be considered to be deactivated when no update is received for three consecutive update intervals. For example, if 192.168.10.1 and 192.168.10.2 do not receive routing table updates from each other for three consecutive times, they are going to consider the link between them is down. Be advised, there might be other ways to detect link failure, but **this is the only method that you are allowed to use to detect link failure**. We can reactivate a link by invoking the **up** command from our driver, as described in subsection 5.2.

6.5 Receiving Messages from the Driver

Each router should be able to receive messages from the driver. As discussed previously, there are two kinds of messages - one for updating link cost and another for passing messages between routers.

6.5.1 Link Update Message

A router should be able to recognize the link update messages as discussed in subsection 5.2.

6.5.2 Send Message

A router should be able to recognize the send messages as discussed in subsection 5.2. Once a router receives the send command from our driver, it should perform necessary packet forwarding activities.

6.6 Packet Forwarding

We are using the work 'packet' in a pretty loose sense here, any message between two routers is a 'packet'. Suppose we invoke the send command as follows:

```
send 192.168.10.1 192.168.10.4 5 hello
```

Let us further assume that this hello packet takes the following path to reach from 192.168.10.1 to 192.168.10.4:

```
192.168.10.1 -> 192.168.10.3 -> 192.168.10.2 -> 192.168.10.4
```

Actually this is the exact same path the packet is supposed to take considering our given topology. Now once 192.168.10.1 receives the send command, it should lookup its routing table to find the next hop to reach 192.168.10.4. It's

going to find that the next hop is 192.168.10.3, so it should send a message to 192.168.10.3:

```
frwd 192.168.10.4 5 hello
```

Be advised, **this message should be sent to 192.168.10.3 (the next hop router) only**. Once 192.168.10.3 receives the forward message, it in turn performs a table lookup and forwards the packet. This goes on until the packet reaches the destination. **You need to print messages on the console stating these forwarding events**. So a typical sequence of messages for the above scenario might be something like this:

```
hello packet forwarded to 192.168.10.3 (printed by 192.168.10.1)
hello packet forwarded to 192.168.10.2 (printed by 192.168.10.3)
hello packet forwarded to 192.168.10.4 (printed by 192.168.10.2)
hello packet reached destination (printed by 192.168.10.4)
```

7   A Sample Sequence of Driver Commands

In the previous sections, we took a look at all the available driver commands and their possible effects. Just for the sake of cohesiveness, let's see the behavior of our routers for a sequence of driver commands. We are going to work on the same topology that we have been studying so far, see Figure 1 to jog your memory. Recall that you can instruct a router to show its routing table by the `show <ip>` command.

1. **Run all the routers, then run the driver**.
   The routers show their initial routing table as each of them are run.

2. **send 192.168.10.1 192.168.10.4 5 hello**
   'hello' reaches destination following this path: 1 -⟩ 3 -⟩ 2 -⟩ 4.

3. **cost 192.168.10.1 192.168.10.2 2**
   Routing tables are updated as required.

4. **send 192.168.10.1 192.168.10.4 11 hello again**
   'hello again' reaches destination following this path: 1 -⟩ 2 -⟩ 4.

5. **down 192.168.10.1 192.168.10.2**
   OS level call made to drop packets through the link between 1 and 2.

6. **send 192.168.10.1 192.168.10.4 3 bye**
   Assuming the above command was invoked before the adjacent routers detected link failure and converged, this message will not reach destination.

7. **send 192.168.10.1 192.168.10.4 9 bye again**
   Assuming the adjacent routers detected the failure by this time and they converged, 'bye again' will reach destination following this path: 1 -⟩ 3 -⟩ 2 -⟩ 4.

8. If you invoke up command and send the previous message, the message will reach destination using the former path i.e. 1 -⟩ 2 -⟩ 4.

8   Marks Distribution

Tasks	Marks
Construction of the initial routing table and showing it	20
Routing table update on cost change	20
Routing table update on link update	20
Packet forwarding	20
Coding style, use of correct data structures, memory management	20
Total	**100**
Bonus	**10**